



THE UNIVERSITY OF
SYDNEY

BUSINESS THESIS

DISCIPLINE OF BUSINESS ANALYTICS

**NFM-GNN: GRAPH-AUGMENTED NEURAL
FOURIER MACHINES FOR VOLATILITY
FORECASTING**

Author: **Yuqing Zhou**

SID: 530056885

Supervisor: **Prof. Junbin GAO**

*A thesis submitted in fulfillment of the requirements for the courses of
BUSS7001 (Business Thesis)*

**The University of Sydney
BUSINESS SCHOOL**

November 2025

Contents

1	Introduction	5
1.1	Background	5
1.2	Contributions	9
1.3	Thesis Structure	10
2	Literature Review	11
2.1	Econometric Models	11
2.2	Machine Learning Approaches	13
2.3	Frequency-Domain Neural Networks	14
2.4	Graph Neural Networks	15
2.5	Motivation and Research Gap	16
3	The Research Questions	19
3.1	Research Questions	19
3.2	Hypothesis Development	20
4	Methodology	22

4.1	Architectural Overview	22
4.2	Model	23
4.3	Modeling Pipeline	24
5	Data Collection, Feature Engineering and Experiment Settings	28
5.1	Data Source and Description	28
5.2	Feature Engineering and Preprocessing	29
5.3	Experimental Setup	30
5.3.1	Forecasting Task Definition	30
5.3.2	Data Splitting Strategy	31
5.3.3	Evaluation Protocol	31
5.4	Baseline Models	32
6	Results and Interpretation	34
6.1	NFM-GNN Prediction Results	34
6.2	MAE and RMSE Comparison	36
6.3	Ablation Study	38
6.4	Statistical Comparison via DM Test and MCS Test	41
7	Discussion	44
7.1	Interpretation of the Main Findings	44
7.2	Comparison with Existing Studies	45

7.3	Practical and Economic Implications	46
7.4	Limitations and Data Challenges	46
7.5	Directions for Future Research	47
7.6	Integrative Reflection	48

List of Figures

4.1	Illustration of NFM-GNN architecture consisting of four main stages . . .	23
5.1	Cross-market correlation matrix of realized volatility among eight major stock indices (2006–2022). Color intensity represents the degree of interdependence, with strong linkages observed among European markets and between Europe and the U.S. (S&P 500), while Asian indices show relatively weaker correlations.	29
6.1	Predicted vs. actual index values across eight major stock indices (SPX, GDAXI, FCHI, FTSE, OMXSPI, N225, KS11, HSI). The close alignment between predicted and actual trajectories demonstrates the effectiveness of the NFM-GNN framework in capturing both short-term fluctuations and long-term volatility dynamics.	35
6.2	MAE Performance of the Proposed Model Across Eight Market Indices.	36
6.3	RMSE Performance of the Proposed Model Across Eight Market Indices .	36
6.4	Significant wins and losses of NFM-GNN relative to baseline models under MAE (left) and RMSE (right) criteria at the 5% significance level based on Diebold–Mariano (DM) tests.	42

List of Tables

6.1	Volatility Forecasting Performance (MAE)	37
6.2	Volatility Forecasting Performance (RMSE)	37
6.3	MAE Comparison across Variants	40
6.4	RMSE Comparison across Variants	40
6.5	Ablation Study: Macro-Averaged Relative Change ($\overline{\Delta}$) in MAE and RMSE	41
6.6	MCS Test (MAE)	43
6.7	MCS Test (RMSE)	43

Abstract

Volatility forecasting is a central challenge in financial modeling, particularly under non-stationary and rapidly evolving market conditions. Existing models often assume feature independence or focus solely on global temporal dynamics, overlooking the complex and transient relationships that emerge across time, features, and assets. While spectral-based models such as Neural Fourier Machines (NFM) achieve strong performance in multi-hundred-step forecasting horizons, they struggle to capture 22-day volatility, an important monthly scale in financial time-series applications. To address this limitation, this paper proposes NFM-GNN, a hybrid model that integrates graph-based attention modules with frequency-aware forecasting. It introduces three graph structures: a temporal graph that captures volatility propagation over time, a feature graph that encodes interactions within each stock index, and a spatial graph that models dynamic co-movement among financial indicators. By preserving the three-dimensional structure of financial time series and performing structure-aware encoding before spectral modeling, NFM-GNN effectively bridges local adaptivity and global generalization. Empirical results on eight representative stock indices show significant improvements in 22-step-ahead forecasts. We further validate the NFM-GNN model through ablation experiments, Model Confidence Set analysis, and Diebold–Mariano tests.

Chapter 1

Introduction

1.1 Background

The stock market has long served as one of the most prominent asset classes, where investors allocate capital in pursuit of higher returns than those offered by risk-free or low-yield assets such as government deposits (Sharpe, 1964). However, this pursuit inherently involves greater exposure to uncertainty and risk, as financial markets are driven by evolving information, expectations, and risk premia (Fama, 1991; Merton, 1973). Financial markets are inherently complex, influenced by dynamic internal and external economic conditions, geopolitical developments, and technological evolution (Cont, 2001). In this context, *volatility* has emerged as one of the most fundamental indicators of market risk and instability.

Volatility forecasting is a cornerstone of modern financial decision-making. Volatility reflects the amplitude and frequency of price fluctuations, embodying the degree to which information, expectations, and shocks are absorbed into asset prices. Accurate forecasts inform key activities such as portfolio allocation, risk management, derivative pricing, and regulatory capital planning (Poon and Granger, 2003). For institutional investors, volatility estimates guide strategic asset allocation and hedging decisions; for algorithmic

traders, they underpin position sizing and real-time risk control. At the macro level, central banks and regulators rely on volatility measures to monitor systemic risk and calibrate macroprudential policies, while recent evidence suggests that volatility-based indicators may also function as early-warning signals of financial bubbles (Kürüm, Weber, and Iyigun, 2018), underscoring their diagnostic value for policymakers and their central role in safeguarding financial stability.

Early empirical research established the importance of volatility modeling in both academic work and financial practice. Andersen and Bollerslev (1998) demonstrated that GARCH models could produce forecasts that closely tracked realized variance, providing one of the first convincing pieces of evidence that volatility could be meaningfully measured and predicted. This empirical progress motivated the formal definition of realized volatility (RV) by Barndorff-Nielsen and Shephard (2002), who developed its econometric foundations, verified its desirable statistical properties, and linked it to continuous-time stochastic volatility models. By allowing volatility to be directly observed from high-frequency return variation rather than inferred from model assumptions, RV marked a paradigm shift in financial econometrics. It offered a data-driven measure of true market variance, bridging theoretical modeling and empirical observation, and has since become a benchmark indicator for both research and practice.

Building on this foundation, Andersen, Bollerslev, Christoffersen, and Diebold (2007) showed how RV-based models could be employed in real-world risk management, including portfolio-level monitoring, asset-level forecasting, and regulatory capital planning. Their work effectively bridged academic theory and institutional application, making realized volatility an operational tool for investors and regulators alike.

Taken together, these developments underscore why volatility forecasting matters: it serves as the analytical bridge between econometric theory, financial practice, and regulatory policy. Reliable forecasts enhance derivative pricing, optimize portfolios, and strengthen systemic risk assessment—making volatility not merely a statistical measure, but a vital lens through which financial stability is understood and managed.

Despite these advances, volatility series remain notoriously difficult to forecast. They are highly nonlinear, noisy, and often dominated by irregular bursts rather than smooth trends (L. Chen, Pelger, and J. Zhu, 2024). Empirical evidence shows that oil prices, policy uncertainty, and investor sentiment exert significant effects on market volatility, while the roles of institutional fund flows and algorithmic trading remain debated and inconclusive (Dhingra, Batra, Aggarwal, Yadav, and Kumar, 2024). Moreover, volatility rarely stays contained within a single market. Through financial linkages and correlated investor behavior, shocks originating in one region can quickly propagate to others, a phenomenon widely known as the volatility spillover effect. Studies have documented such interconnected dynamics between China and Southeast Asian stock markets (Zhong and J. Liu, 2021), and highlighted the amplifying influence of geopolitical risk, policy uncertainty, and climate risk across economic cycles (Z. Hu and Borjigin, 2024). Such interconnected market behavior suggests that volatility is not merely a time-series phenomenon but part of a broader relational structure.

In this sense, the ability to model and predict volatility with high fidelity is not simply an academic exercise but a pressing necessity for the stability and efficiency of global financial markets. This raises a central question: how should realized volatility (RV) be forecast effectively in settings that are noisy, nonlinear, and structurally dynamic?

Traditional econometric approaches, most notably the ARCH (Engle, 1982) and GARCH (Bollerslev, 1986) families, laid the groundwork by capturing conditional heteroskedasticity. Extensions such as EGARCH (D. B. Nelson, 1991), GJR-GARCH (Glosten, Jagannathan, and Runkle, 1993), and the heterogeneous autoregressive (HAR) model (Corsi, 2009) further address asymmetries and multi-horizon persistence. A research on lassoing the HAR model (Aldrino and Knaus, 2016) applied the Lasso method to the HAR framework, showing that while HAR remains parsimonious and robust, its lag structure may deviate across assets and reveal structural breaks. While theoretically elegant and widely applied, these models rely on strong distributional and linearity assumptions. They struggle with nonlinearities, structural breaks, and cross-series interactions—features that intensify during market turbulence—so their predictive performance can deteriorate under

regime shifts.

Learning-based methods attempt to relax these constraints (Christensen, Siggaard, and Veliyev, 2023). Early machine-learning models like Random Forests (Luong and Dokuchaev, 2018)) reduced parametric reliance but depended on manual feature engineering. Deep neural architectures—ranging from recurrent models (e.g., LSTMs) and CNN–LSTM hybrids to decomposition-based forecasters like N-BEATS (Souto and Moradi, 2024) and NHITS (Souto, 2025)—have improved the modeling of complex temporal patterns via data-driven representations. In parallel, frequency-domain approaches such as FED-former (Zhou, Ma, Wen, Xue Wang, Sun, and R. Jin, 2022) and Neural Fourier Machines (NFM) (Kim, Hioka, and Witbrock, 2024) demonstrate that spectral representations help capture long-range dependencies efficiently. Despite these advances, a common limitation persists: most models treat volatility features as independent sequences. This simplification overlooks the inherently relational nature of financial volatility.

In reality, volatility spillovers and co-movements across assets form a graph-like structure, where nodes represent markets or indices and edges capture evolving dependencies such as correlations, information flows, or shared shocks. These relationships are dynamic with changing macroeconomic regimes and market sentiment. They span multiple temporal scales (Yanan Li and Giles, 2015; Tsai, 2014; Cheng and Q. Li, 2021). Short-horizon volatility often reflects microstructural noise and trading activity, whereas long-horizon trends embody macroeconomic cycles and systemic expectations. Ignoring these structured interactions risks losing information essential for understanding how localized shocks propagate through the system and reshape aggregate risk.

Based on these limitations, we reconsider how volatility should be represented and predicted. While time-domain models can capture sequential variations directly, financial volatility signals are often dominated by irregular bursts and noise, making their temporal patterns unstable and difficult to model. Recasting the problem in the frequency domain allows these fluctuations to be interpreted as combinations of periodic components, thereby filtering out transient shocks and revealing the underlying cyclical structure of

market behavior. This perspective mitigates the randomness of large, short-lived swings and provides a smoother representation for forecasting. However, merely operating in the frequency domain may overlook the structured interactions that shape volatility dynamics. An effective volatility model should therefore combine spectral representation with structure-aware learning, leveraging graph-based formulations and domain-informed priors (e.g., moving averages, sectoral linkages) to jointly capture cyclical patterns, relational dependencies, and cross-scale interactions.

1.2 Contributions

Motivated by these research gaps, we integrate structure-aware graph learning with frequency-based forecasting. Our model builds relational graphs along three axes (i.e., time, features and entities) and incorporates them into a Fourier-based forecasting backbone. Our contributions are as follows:

- We identify a structural gap in frequency-based volatility modeling: the lack of axis-wise relational encoding across time and features.
- We propose a novel architecture, NFM-GNN, that integrates three graph-based attention modules prior to Fourier forecasting.
- We demonstrate the effectiveness of our approach through comprehensive evaluations, showing strong gains in long-horizon robustness.

By embedding these relational graphs into a Fourier-based forecasting backbone, our framework seeks to improve both the robustness of long-horizon forecasts and the adaptability to short-term volatility spikes.

1.3 Thesis Structure

The remainder of this article is structured as follows. Chapter 2 provides a review of the literature on econometric and machine learning approaches to volatility forecasting. Chapter 3 sets out the research question. Chapter 4 introduces the proposed framework, emphasizing the integration of graph-based relational modules with Fourier-based forecasting. Chapter 5 details the data sources, the construction of the realized volatility, and the pre-processing procedures. Chapter 6 presents the empirical results together with statistical evaluations of predictive precision and offers an interpretation of the findings. Finally, Chapter 7 discusses the broader implications for academic research and financial practice and concludes with suggestions for future research.

Chapter 2

Literature Review

We structure the review along four strands: (i) econometric volatility models, (ii) machine learning, (iii) frequency-domain architectures, and (iv) graph-based approaches. Each strand is critically examined in terms of its strengths and limitations, highlighting how they have contributed to but also constrained progress in volatility forecasting. Building on this synthesis, the review identifies the remaining research gaps and establishes the motivation for the proposed framework presented in this thesis.

2.1 Econometric Models

Classical econometric models such as ARCH (Engle, 1982) and GARCH (Bollerslev, 1986) established the foundation of volatility analysis by modeling conditional heteroskedasticity. Engle (1982) introduced the autoregressive conditional heteroskedastic (ARCH) process to relax the implausible assumption of a constant one-period forecast variance. ARCH models define mean-zero, serially uncorrelated processes with nonconstant conditional variances but constant unconditional variances, allowing recent past shocks to inform the one-step-ahead variance. Bollerslev (1986) subsequently generalized this framework into the GARCH model, in which current conditional variance depends not only on past squared disturbances but also on its own lagged values. The paper derives stationarity

conditions and autocorrelation structures for this parametric class and discusses estimation and testing via maximum likelihood. However, despite their empirical success, both ARCH and GARCH models exhibit notable shortcomings. They assume a symmetric response of volatility to positive and negative shocks. Such simplifications limit their ability to reproduce the leverage effect and the heterogeneous temporal dynamics observed in real markets.

To address these issues, a series of extensions have been proposed. Their extensions, including EGARCH (D. B. Nelson, 1991) and GJR-GARCH (Glosten, Jagannathan, and Runkle, 1993), incorporated leverage and asymmetry effects to better capture stylized facts of financial markets. Empirical findings by Müller, Dacorogna, Davé, Olsen, Pictet, and Von Weizsäcker (1997) revealed that market participants operating at heterogeneous time horizons generate volatility components that interact across temporal scales, producing long-memory behavior and cross-scale predictability. Motivated by this evidence, the heterogeneous autoregressive (HAR) model (Corsi, 2009; Corsi, Audrino, and Renò, 2012) further enriched this literature by incorporating realized volatility across daily, weekly, and monthly horizons, offering a multi-scale perspective. It effectively reproduces key empirical features of financial returns, including long memory, fat tails, and self-similarity, while maintaining parsimony and tractability. However, its additive linear structure hinders the modeling of nonlinearities, breaks, and regime changes.

While influential, these econometric approaches remain heavily assumption-driven, relying on linear additivity and stationarity. As a result, they often underperform in the presence of nonlinearities, structural breaks, and cross-asset volatility spillovers. These limitations highlight the need for more flexible approaches that can adapt to regime shifts and capture multivariate dependencies.

2.2 Machine Learning Approaches

Machine learning (ML) methods emerged as an alternative to address the rigidity of econometric models. Algorithms such as Support Vector Regression (SVR) (Vapnik, 1995), Random Forests (Breiman, 2001), and XGBoost (T. Chen and Guestrin, 2016) demonstrated improved forecasting accuracy by capturing nonlinear dependencies without strict distributional assumptions. Nevertheless, their reliance on hand-crafted lagged features constrains scalability and transferability, and they lack explicit mechanisms to capture temporal dynamics.

Recent models extended this paradigm by enabling automatic representation learning and long-horizon modeling. Recurrent networks and LSTMs (Hochreiter and Schmidhuber, 1997; D. M. Nelson, Pereira, and Oliveira, 2017) effectively capture sequential dependencies, while CNN–LSTM hybrids (Mehtab, Sen, and Dasgupta, 2020) combine local and global temporal structures. But they often struggle with sudden changes in temporal patterns and complex feature interactions.

Basis-expansion models have been proposed: N-BEATS (Boris N Oreshkin, Carпов, Chapados, and Bengio, 2019), a deep architecture with residual links and fully connected layers, is interpretable, domain-agnostic, and computationally efficient; Its extension, NBEATSx (K. G. Olivares, Challu, Marcjasz, Weron, and Dubrawski, 2023), further projects the target signal onto nonlinear basis functions and incorporates exogenous variables, improving forecasting accuracy; The NHITS model (Challu, S. G. Olivares, Boris N. Oreshkin, Ramirez, Canseco, and Dubrawski, 2023), as a generalization and extension of N-BEATS, further enhances hierarchical forecasting through multi-rate data sampling and hierarchical interpolation, achieving strong performance across long-horizon forecasting tasks. However, they lack a mechanism for modeling interactions between variables to adequately capture cross-feature or cross-asset dependency structures in financial series.

Despite their advances, most deep learning methods are trained directly in the time do-

main. This design makes them vulnerable to distributional shifts and limits their ability to explicitly model cross-asset linkages, thereby reducing robustness in periods of financial turbulence.

Taken together, ML have substantially advanced volatility forecasting by improving flexibility and adaptability beyond econometric baselines. Yet, their limitations, manual feature dependence in ML, and time-domain sensitivity with weak cross-sectional modeling, underscore the need for alternative representation domains and structured models that can integrate both temporal and cross-asset dynamics.

2.3 Frequency-Domain Neural Networks

An alternative stream of research leverages frequency-domain representations. Autoformer (H. Wu, J. Xu, J. Wang, and Long, 2021), FEDformer (Zhou, Ma, Wen, Xue Wang, Sun, and R. Jin, 2022), FourierGNN (K. Yi, Q. Zhang, Fan, H. He, L. Hu, P. Wang, An, Cao, and Niu, 2023) and Neural Fourier Machines (NFM) (Kim, Hioka, and Witbrock, 2024) exploit spectral decompositions to model long-range dependencies and periodicities efficiently.

FourierGNN simplifies multivariate time-series forecasting by unifying spatial and temporal dependencies within a single graph-learning framework. Through a Fourier Graph Operator that performs spectral-domain graph convolution, it reduces model complexity while maintaining strong expressive power and efficiency across multiple datasets. However, this formulation is tailored for general MTS tasks and does not preserve the temporal hierarchy and feature semantics crucial for volatility forecasting in financial markets.

NFMs enables strong extrapolation capabilities, robustness to distribution shifts, and efficient long-sequence modeling. NFM reinterprets time series forecasting as a learnable interpolation and extrapolation problem in the frequency domain. Rather than operating directly in the temporal space, it represents input sequences through discrete Fourier trans-

forms (DFT) and performs global sequence transformation via frequency manipulation, enabling it to efficiently model long-range structures and periodic patterns. This design is length-independent, instance-specific, and frequency-adaptive. It acts as a global convolution kernel in the spectral space, leveraging the convolution theorem for efficient global attention. The resulting latent representation $\hat{z}$ is defined over a rescaled spectral domain that matches the target output length L , enabling interpolation or extrapolation from the input x to the output y directly in the frequency domain.

However, NFM excels in long-horizon forecasting scenarios, where the prediction length spans hundreds of time steps. It tends to underperform in under 100 time steps horizon forecasting tasks, particularly those with low forecast horizons. While NFM demonstrates strong extrapolation ability and global pattern modeling, it lacks the ability of learning 22 time steps, fine-grained volatility dynamics commonly observed in financial time series.

2.4 Graph Neural Networks

Graph Neural Networks (GNNs) provide a principled framework for modeling interconnected dependencies among financial assets by representing each asset as a node and encoding correlations, spillovers, or causality links as edges (Feng, X. He, Xiang Wang, Luo, Yiqun Liu, and Chua, 2019; K. Xu, Y. Wu, Xia, Sang, and B. Wang, 2022). By jointly learning both the node representations (asset-level states) and edge weights (relationship intensities), GNNs offer a flexible means to capture volatility spillovers, contagion effects, and systemic risk pathways that static models cannot represent. Empirical evidence shows that GNN-based models can enhance multivariate volatility forecasting by explicitly capturing volatility transmission channels (Son, Y. Lee, Park, and J. Lee, 2023), while dynamic graph construction methods adapt to evolving market structures (Chi, Gao, and C. Wang, 2024).

In related domains, spatiotemporal GNNs have achieved success in traffic forecasting (Yaguang Li, R. Yu, Shahabi, and Yan Liu, 2018; B. Yu, Yin, and Z. Zhu, 2018) and

other high-dimensional systems, underscoring their generality for structured time series. Notably, these studies highlight that GNNs can serve as a bridge between relational structures and sequential dynamics, making them a compelling foundation for hybrid architectures that combine temporal modeling with structural reasoning. While most existing financial GNNs rely on real-valued graph representations, recent progress in adjacent engineering and physics domains demonstrates that advanced GNN formulations can capture richer interaction patterns. They have been applied to array antenna radiation pattern prediction (J. Jin, Su, Y. Xu, Z. He, and Lu, 2022), complex matrix inversion (Q. Yi, Xiao, Y. Zhang, Liao, Ding, and Peng, 2018), and spatiotemporal modeling in power systems (T. Wu, Scaglione, and Arnold, 2024). These successes collectively underline the flexibility of graph-based learning across domains involving interdependent signals and dynamic spatial–temporal relationships.

However, despite this progress, most financial GNN studies remain limited to static cross-asset graphs that overlook dependencies within an individual asset’s feature space or fail to adaptively represent relationships across both time and features. Consequently, they capture only part of the volatility propagation mechanism, leaving intra-series feature interactions largely unexplored.

2.5 Motivation and Research Gap

Across econometrics, machine learning (ML), frequency-domain architectures, and graph neural networks (GNNs), prior research has substantially advanced volatility forecasting but remains fragmented. Econometric models provide interpretability but lack flexibility; ML methods improve adaptability but are prone to regime shifts; frequency-domain approaches enhance robustness and extrapolation but often neglect cross-asset interactions.

Volatility forecasting in financial time series fundamentally requires capturing both temporal dynamics and inter-feature dependencies across multiple indices. Yet, most existing architectures, from recurrent networks to transformer-based models, treat input variables

such as realized volatility (RV) as temporally aligned but mutually independent. This simplifying assumption overlooks the complex, time-varying, and asymmetric relationships among features, especially under non-stationary market conditions.

For example, during market shocks, short-horizon indicators may abruptly dominate volatility dynamics, while in stable periods, long-term signals often carry greater predictive weight. Modeling such transient relevance shifts demands architectures capable of dynamically adjusting attention across both time and features—an ability largely absent in conventional sequential models.

Meanwhile, frequency-based frameworks such as the Neural Fourier Machine (NFM) (Kim, Hioka, and Witbrock, 2024) have shown strong performance in long-range forecasting by learning global spectral representations. However, their inherent bias toward stationary and smooth patterns limits their ability to handle short-horizon scenarios characterized by volatility clustering, abrupt regime transitions, and localized feature entanglement.

These limitations reveal a clear research gap: the need to combine frequency-domain representation learning with graph-based relational reasoning. Such integration enables the model to jointly capture periodic patterns and cross-asset structural dependencies, improving robustness and adaptability across heterogeneous markets. Although this direction remains unexplored in financial econometrics, related progress in signal processing and time-series modeling highlights its strong potential.

Furthermore, to the best of our knowledge, no prior work employs GNNs to construct temporal graphs within a single asset’s time series. Existing GNN-based frameworks primarily focus on cross-asset spillovers or market co-movements, treating each asset as a static node. This paradigm overlooks three critical aspects essential for accurate volatility forecasting:

- (1) Fine-grained temporal propagation within each asset;
- (2) Structured relationships among financial indicators;

(3) Evolving cross-asset interactions over time.

Motivated by these gaps, the present study develops an NFM–GNN architecture that integrates frequency-based temporal modeling with graph-based structural learning to enhance volatility forecasting accuracy and stability across markets.

Chapter 3

The Research Questions

3.1 Research Questions

Volatility prediction remains a central yet unresolved challenge in financial econometrics and data science. Existing approaches provide valuable insights but remain fragmented. Motivated by those gaps in 2.5, this thesis addresses the following research questions:

RQ1. Unified Time–Frequency Representation. How can spectral representations derived from Fourier transforms be integrated with graph-based temporal dependencies to jointly capture long-term cyclical trends and abrupt volatility shifts, while preserving their mutual interactions under non-stationary market regimes?

Rationale: This question extends beyond isolated modeling of smooth or burst-like behaviors, aiming to reconcile global periodic structures with local regime changes, a central difficulty in representing financial instability.

RQ2. Axis-Wise Structural Integration. Which relational structures most effectively express the latent organization of financial systems, and how can multi-axis graph representations guide spectral learning to encode these interdependencies?

Rationale: This question explores the deeper geometry of financial data, treating

markets not as independent sequences but as interacting systems with dynamic connectivity.

RQ3. Hybrid Graph Priors and Domain-Informed Learning. How can statistical priors (e.g., correlation matrices, moving averages) be fused with adaptive attention-based graphs to balance global stability and local adaptability in modeling volatility propagation?

Rationale: This examines how domain-informed priors can act as anchors that regularize learned relationships, preventing overfitting while preserving responsiveness to emerging market structures.

RQ4. Dynamic Regime Adaptation. Can a unified spectral–graph framework dynamically adapt to structural disruptions, such as crisis periods, asynchronous trading, and changing forecast horizons and reveal how volatility regimes evolve and reorganize across markets?

Rationale: Beyond robustness testing, this investigates whether the model can uncover regime-dependent relational shifts, providing insight into the systemic fragility.

3.2 Hypothesis Development

Building upon the research questions formulated in Chapter 3.1, this chapter develops empirically testable hypotheses derived from theoretical and empirical insights discussed in the literature review. These hypotheses guide the model design and experimental validation presented in Chapter 4.

H1. Frequency–Temporal Synergy Integrating spectral representations with temporal graph dependencies significantly improves short-horizon volatility forecasting accuracy compared to purely time-domain or frequency-domain models.

H2. Multi-Axis Structural Learning Joint learning across temporal, feature, and spatial (cross-asset) graph axes yields superior forecasting performance by preserving the intrinsic multi-dimensional organization of financial volatility.

H3. Hybrid Graph Priors Models that integrate statistical priors (e.g., correlation matrices, moving averages) with attention-based dynamic adjacency achieve higher stability and lower overfitting risk than models using learned or fixed graphs alone.

H4. Regime Robustness The proposed spectral–graph framework maintains forecasting robustness and relational interpretability under structural disruptions (e.g., COVID-19, asynchronous trading) compared with econometric and deep learning baselines.

Chapter 4

Methodology

4.1 Architectural Overview

Our model, NFM-GNN, introduces structure-aware encoding into the frequency-based forecasting pipeline by decomposing the architecture into four clearly defined stages, as illustrated in Figure 4.1. In Stage 1, the input data are processed by three graph neural networks (GNNs) that respectively capture temporal, feature-level, and spatial dependencies. In Stage 2, their outputs undergo cross-projection and dynamic fusion via an MLP, where each GNN output passes through both a linear and a sinusoidal (periodic) projection, resulting in six representations that are adaptively fused through weighted aggregation. In Stage 3, the fused sequence is extended and encoded by the Learnable Frequency Token (LFT) module, which combines zero-padding with learnable spectral priors and performs channel mixing through an MLP and global convolution to form predictive embeddings. Finally, in Stage 4, an MLP mixer and prediction head generate the final forecasting sequence of length H .




Figure 4.1: Illustration of NFM-GNN architecture consisting of four main stages

4.2 Model

Before we introduce our proposed NFM-GNN model, we define the following notation. Let V be the number of financial indices, T the number of historical time steps, and D the number of features per index (e.g., RV, 5MA, 22MA). At each time step t , the structured input matrix is $\mathbf{X}_t \in \mathbb{R}^{V \times D}$. Stacking over time yields a 3D tensor:

$$\mathcal{X} = [\mathbf{X}_1, \mathbf{X}_2, \dots, \mathbf{X}_T] \in \mathbb{R}^{T \times V \times D}. \quad (4.1)$$

Unlike conventional architectures that flatten inputs along feature axes, we preserve this three-dimensional structure to maintain semantic separation: Temporal axis T : evolution across time; Spatial axis V codes index-wise separation; and Feature axis D statistical indicators. This structure facilitates axis-specific attention and avoids premature fusion of heterogeneous signals.

Our model is a structured hybrid architecture that combines local graph-based relational encoding with global spectral modeling. It operates directly on the 3D input tensor $\mathcal{X}$ defined in Eq. (4.1), and makes prediction for $\mathcal{Y} \in \mathbb{R}^{H \times V \times D}$, where H is forecasting horizon.

4.3 Modeling Pipeline

Stage 1: Axis-wise graph encoding. To explicitly capture intra-series and intra-feature dynamics, we apply three axis-specific graph modules to the raw input $\mathcal{X}$. The temporal graph module (GAC) treats the T time steps of each index as nodes and learns graph-based dependencies across time, producing a sequence-level attention-representation $\mathbf{X}_{\text{tg}} \in \mathbb{R}^{T \times V \times d}$ where d is the hidden dimension used in the attention layers. The process can be defined as

$$\mathbf{X}_{\text{tg}} = \text{att}_{tg}(\mathcal{X}).$$

where $\text{att}(\cdot)$ means attention transform with relevant Q, K, V parameters.

The feature graph module models within the assets between D features within each time step and index, presented as

$$\mathbf{X}_{\text{feat}} = \text{att}_{feat}(\mathcal{X}) \odot |\text{corr}(\mathcal{X})|.$$

Here, $\odot$ denotes element-wise multiplication. The correlation adjacency is derived from the Pearson correlation matrix computed over normalized features. We apply correlation-based modulation only to the feature graph because feature dimensions (e.g., RV5, 5MA, 22MA) are semantically aligned and often exhibit strong statistical dependencies, while temporal and asset graphs are more context-dependent and dynamic, where hard-coded correlation may bias or oversimplify complex patterns.

The spatial graph module models co-movements among the assets between V features within each time step and features, yielding $\mathbf{X}_{\text{asset}}$ of the same shape, presented as

$$\mathbf{X}_{\text{asset}} = \text{att}_{asset}(\mathcal{X}).$$

These three outputs capture localized temporal propagation and feature-wise entanglement under nonstationary dynamics, by exploiting structure information along temporal,

feature, and indices axes through relevant attention graphs. This node-based interpretation provides a principled view of attention as a soft graph construction process. Instead of relying on fixed graph structures or external priors, NFM-GNN model learns data-specific, context-dependent relational structures tailored for each time-index pair.

Remark: A temporal attention graph explicitly models relational dependencies between time steps for each financial index, allowing the model to go beyond simple autoregressive assumptions. Financial volatility exhibits irregular patterns, delayed effects, and *non-Markovian dependencies*, where, for example, a shock several days ago might still influence today’s volatility due to latent structural persistence. By treating each time step as a node in a temporal graph and allowing flexible graph-based connections, the model can learn to identify such non-local and heterogeneous dependencies. This includes **Short-term bursts**: capturing volatility clustering over a few steps; **Mid-term echoes**: detecting periodic re-emergence of volatility patterns; and **Asymmetric memory**: modeling how certain patterns decay slower or faster over time.

Stage 2: Cross-projection and latent fusion. Each of the three graph outputs is passed through two projection branches. Let $\mathbf{X}^{(1)} = \mathbf{X}_{\text{tg}}$, $\mathbf{X}^{(2)} = \mathbf{X}_{\text{feat}}$, and $\mathbf{X}^{(3)} = \mathbf{X}_{\text{asset}}$, with each $\mathbf{X}^{(i)} \in \mathbb{R}^{T \times V \times d}$. We apply both projections to each $\mathbf{X}^{(i)}$, resulting in six projected tensors. The first, proj_1 , is a shared linear projection that performs a standard linear transformation to preserve direct feature relationships:

$$\text{proj}_1(\mathbf{X}^{(i)}) = \mathbf{W}_1 \mathbf{X}^{(i)}.$$

The second, proj_2 , is a nonlinear periodic projection based on sinusoidal activations, which introduces spectral nonlinearity by expanding the feature dimension, applying sinusoidal modulation, and then compressing the result back to the hidden dimension:

$$\text{proj}_2(\mathbf{X}^{(i)}) = \mathbf{W}_3 \phi_{\text{mix}}(\omega \mathbf{W}_2 \mathbf{X}^{(i)}),$$

where $\mathbf{W}_1, \mathbf{W}_2, \mathbf{W}_3$ are learnable weight matrices, $\phi_{\text{mix}}(\cdot)$ is a periodic activation func-

tion combining sinusoidal components (e.g., $\sin(x) + \cos(x)$), and ω is a frequency coefficient controlling the scale of sinusoidal oscillation.

This results in six outputs and through an Multi-Layer Perceptron these are summed weightedly. A fusion weight vector $\mathbf{w} \in \mathbb{R}^6$ is dynamically generated via an MLP from global context statistics. The final fused representation $\bar{\mathbf{X}} \in \mathbb{R}^{T \times V \times D}$ is computed as in 3 axes:

$$\bar{\mathbf{X}} = \sum_{i=1}^3 \left(w_{2(i-1)+1} \cdot \mathbf{proj}_1(\mathbf{X}^{(i)}) + w_{2(i-1)+2} \cdot \mathbf{proj}_2(\mathbf{X}^{(i)}) \right).$$

This fused latent sequence $\bar{\mathbf{X}}$ encodes both graph-informed structure and frequency-enhanced expressiveness, and serves as the input to the spectral modeling stage.

Stage 3: Spectral modeling via LFT and INFF.

To produce the forecasting sequence of length H , the fused sequence $\mathbf{X}$ is passed to a Learnable Frequency Token (LFT) module that generates an overall sequence of length $T + H$ by zero-padded extension plus learnable spectral priors.

First, we conduct Fourier transform $\mathbf{X}_{\text{freq}} = FFT(\mathbf{X})$ over the time axis, resulting Fourier signal $\{\mathbf{X}_{\text{freq}}[k]\}_{k=0}^{T-1}$ of length T . Then $\{\mathbf{X}_{\text{freq}}[k]\}$ will be scaled into length $T + H$ by interpolating at linearly scaled frequency positions with an implicit neural representation (SIREN-style) (Sitzmann, Martel, Bergman, Lindell, and Wetzstein, 2020). For example, in the case of $H = T$, a new frequency component will be added between $\mathbf{X}_{\text{freq}}[k]$ and $\mathbf{X}_{\text{freq}}[k + 1]$.

This yields $\mathbf{Z}_0$ of length $T + H$, a length-adaptive spectral embedding. We then apply a stack of Mixer blocks, following the Implicit Neural Fourier Filter (INFF) design (Kim, Hioka, and Witbrock, 2024). Each Mixer layer performs token mixing in the frequency domain, conditioned on $\mathbf{Z}_0$, followed by INFF which conducts Fourier filtering $\mathbf{Z}[k] = W[k] * \mathbf{Z}_0[k]$. The learnable Fourier $\{W[k]\}_{k=1}^{T+H-1}$ are complex-valued and generated by another SIREN function defined on frequency index k . All these transform are conducted along the frequency domain independently of spatial dimension and feature dimension.

The output $\mathbf{z}$ is refined through residual connections, normalization layers and a final feed-forward layer, producing the final global latent representation $\mathbf{Z} \in \mathbb{R}^{(T+H) \times V \times d}$ by the inverse Fourier transform with the last H steps $\hat{\mathbf{Y}} = \mathbf{Z}[-H :]$ as model prediction.

Stage 4: Downstream prediction. The last H slices of the latent $\mathbf{z}$ from Stage 3 can be projected from d latent features into the appropriate feature dimension in any task-specific prediction head. In our volatility forecasting setup, it supports horizon-adaptive forecasting while retaining both local structural and global periodic information.

Overall, the model introduces structured graph inductive bias before entering the frequency space, ensuring that LFT and INFF modules operate over feature-aware, temporally grounded representations. This design enables the system to handle short-horizon, non-stationary volatility patterns with stronger structural adaptivity than pure spectral models.

The Training Loss Function. The training objective combines time-domain prediction accuracy with frequency-domain consistency, ensuring that the model learns both short-term fluctuations and long-range spectral coherence. Empirically, we find that incorporating frequency-domain supervision substantially improves the stability of training and the fidelity of periodic representations. The loss function is defined as

$$\mathcal{L} = w_e \|\hat{\mathcal{Y}} - \mathcal{Y}\|^2 + w_f \|\hat{\mathbf{Y}}_{\text{freq}} - \mathbf{Y}_{\text{freq}}\|^2.$$

where $w_e > 0$ and $w_f > 0$ are weighting coefficients that balance errors across the time and frequency domains. Unlike a simple equal weighting (e.g., $w_e = w_f = 0.5$), these coefficients are empirically tuned to account for the different energy scales and noise sensitivities between the two domains. The frequency-domain term helps constrain the model to preserve cyclical structures, while the time-domain term ensures alignment with observed short-horizon variations. Both $\hat{\mathbf{Y}}_{\text{freq}}$ and $\mathbf{Y}_{\text{freq}}$ are Fourier transforms of model output $\mathbf{z}$ and the observed sequence $(\mathcal{X}, \mathcal{Y})$ in length $T + H$.

Chapter 5

Data Collection, Feature Engineering and Experiment Settings

5.1 Data Source and Description

This study utilizes high-frequency financial data sourced from the **Oxford-Man Institute’s Realized Library (v0.4)**, a widely recognized benchmark dataset for volatility research. Our analysis covers eight major international stock market indices, selected to represent a geographically diverse and economically significant set of global markets:

- **North America:** S&P 500 (SPX)
- **Europe:** DAX (GDAXI), CAC 40 (FCHI), FTSE 100 (FTSE), OMX Stockholm 30 (OMXSPI)
- **Asia:** Nikkei 225 (N225), KOSPI (KS11), Hang Seng Index (HSI)

The dataset spans from **November 22, 2006, to March 30, 2022**, encompassing a rich variety of market conditions, including the 2008 Global Financial Crisis, periods of relative calm, and the COVID-19 pandemic-induced volatility surge. This extensive timeframe allows for a robust evaluation of the model’s performance across different market regimes.

To visualize the degree of interconnectedness across markets, we compute and plot the correlation matrix of realized volatility among these indices over the full sample period.




Figure 5.1: Cross-market correlation matrix of realized volatility among eight major stock indices (2006–2022). Color intensity represents the degree of interdependence, with strong linkages observed among European markets and between Europe and the U.S. (S&P 500), while Asian indices show relatively weaker correlations.

5.2 Feature Engineering and Preprocessing

For each index, we construct a set of input features designed to capture volatility dynamics at different time scales. The core feature is the daily RV, calculated as the sum of squared 5-minute intraday log-returns.

To ensure the stationarity of the time series and improve numerical stability for deep learning models, all RV data undergoes a standard preprocessing procedure. Specifically, we apply a **square-root transformation**, which helps to stabilize the variance and make

the data’s distribution more Gaussian. Subsequently, the values are multiplied by 100 to scale them into a numerically favorable range for model training. Any missing data points, typically resulting from market holidays, are handled using a forward-fill imputation method.

In addition to the primary RV feature, we engineer two moving average features to provide the model with explicit trend information:

- **5-day Moving Average (5MA)** of the transformed RV, capturing short-term trends.
- **22-day Moving Average (22MA)** of the transformed RV, capturing longer-term (monthly) trends.

Thus, for each of the $V = 8$ indices, the input at any time step consists of $D = 3$ features.

5.3 Experimental Setup

5.3.1 Forecasting Task Definition

The primary objective is to perform multivariate volatility forecasting. The model is tasked with predicting the next H days of realized volatility for all eight indices simultaneously, given a look-back window of the previous T days of historical data. In line with common practice in financial risk management and portfolio allocation, we set the primary forecast horizon to $H = 22$ trading days to approximate a standard monthly window. The choice of the forecast horizon reflects the conventional assumption of approximately 22 trading sessions per month in major equity markets, making it a natural and interpretable horizon for monthly volatility forecasting, portfolio rebalancing, and regulatory risk reporting cycles.

5.3.2 Data Splitting Strategy

To prevent any look-ahead bias and respect the temporal nature of the data, we employ a strict chronological splitting strategy. The entire dataset of N total trading days is partitioned as follows:

1. **Testing Set:** The final 30% of the data is held out as the test set for final model evaluation.
2. **Training and Validation Sets:** The initial 70% of the data is used for model development. This portion is further split, with the first 85% serving as the training set and the subsequent 15% as the validation set for hyperparameter tuning and early stopping.

5.3.3 Evaluation Protocol

Model performance is evaluated on the test set using a rolling-window forecast methodology. Starting from the beginning of the test set, the model is trained on a window of historical data, makes a 22-day-ahead prediction, and then the window is rolled forward. For our experiments, the step size of the roll is set to 22 days to ensure that each evaluation window is independent. This protocol provides a robust assessment of the model’s practical utility in a simulated real-world forecasting scenario. We evaluate the proposed model on multivariate realized volatility forecasting, using square-root daily realized volatility (RV) data computed from 5-minute intraday returns. The primary prediction horizon is set to $h = 22$ trading days, corresponding to a standard monthly investment window, representing a long-term prediction target.

5.4 Baseline Models

To provide a comprehensive and rigorous evaluation of the proposed NFM-GNN model, we establish a strong set of baseline models. The selection is guided by the principle of including representatives from distinct and influential families of time-series forecasting methodologies: classical econometric models, standard recurrent neural networks, and recent state-of-the-art deep learning architectures. This diverse set allows for a multi-faceted assessment of our model’s capabilities.

Econometric Models: These models are the traditional workhorses in financial econometrics and serve as essential benchmarks for interpretability and performance on standard volatility forecasting tasks.

- **GARCH** (Bollerslev, 1986): A benchmark volatility model that captures the phenomenon of volatility clustering by modeling conditional variance as a function of past shocks and lagged variances.
- **HAR** (Corsi, 2009): It regresses future realized volatility on daily, weekly, and monthly averages of past realized volatility (1-, 5-, and 22-day horizons), capturing long-memory effects in a parsimonious way.

Classic Deep Learning Model: This category represents the standard deep learning approach for sequence modeling before the advent of Transformers and more complex architectures.

- **LSTM** (Hochreiter and Schmidhuber, 1997): It serves as a strong baseline for non-linear time-series forecasting and represents the standard deep learning approach that our model should outperform.

State-of-the-Art Time-Series Architectures: We include two recent, powerful deep learning models that have demonstrated top-tier performance on a wide range of forecasting benchmarks. These models represent the current SOTA that NFM-GNN must compete with to prove its advancement.

- **N-BEATSx** (K. G. Olivares, Challu, Marcjasz, Weron, and Dubrawski, 2023): A deep neural architecture based on backward and forward residual links and a stack of blocks with a specific basis expansion. The 'x' variant extends the original N-BEATS to incorporate exogenous variables, making it a powerful and flexible general-purpose forecaster.
- **N-HITS** (Challu, S. G. Olivares, Boris N. Oreshkin, Ramirez, Canseco, and Dubrawski, 2023): A recent SOTA model that leverages hierarchical interpolation and multi-rate signal processing.

For a fair and rigorous comparison, all baseline models are subjected to the same experimental protocol. Each model is implemented within a consistent framework and its key hyperparameters (e.g., hidden dimension, learning rate, block depth, number of layers) are carefully tuned on the validation set using a systematic grid search. All models are trained on the same preprocessed data and evaluated on the identical test set for the $H = 22$ day-ahead forecasting task.

Chapter 6

Results and Interpretation

6.1 NFM-GNN Prediction Results

Figure 6.1 illustrates the time-series alignment between the predicted and actual values for eight major stock indices. The forecasts closely track realized trajectories, indicating that the NFM-GNN framework is able to capture both transient shocks and persistent volatility dynamics across heterogeneous markets.

Figure 6.2 and Figure 6.3 provides a quantitative evaluation based on MAE and RMSE. The results show that prediction errors remain consistently low across all indices. In particular, OMXSPI exhibits the lowest error levels ($\text{MAE} = 0.208$, $\text{RMSE} = 0.376$), while even highly volatile markets such as SPX maintain moderate error magnitudes ($\text{MAE} \approx 0.322$, $\text{RMSE} \approx 0.598$).



Figure 6.1: Predicted vs. actual index values across eight major stock indices (SPX, GDAXI, FCHI, FTSE, OMXSPI, N225, KS11, HSI). The close alignment between predicted and actual trajectories demonstrates the effectiveness of the NFM-GNN framework in capturing both short-term fluctuations and long-term volatility dynamics.




Figure 6.2: MAE Performance of the Proposed Model Across Eight Market Indices.




Figure 6.3: RMSE Performance of the Proposed Model Across Eight Market Indices

6.2 MAE and RMSE Comparison

Tables 6.1 and 6.2 summarize the forecasting accuracy of all models across the eight stock indices. The performance is assessed using both MAE and RMSE to capture overall accuracy and sensitivity to large deviations. Bolded values indicate the best-performing model for each index. Gray-shaded cells denote models that are included in the MCS test

Table 6.1: Volatility Forecasting Performance (MAE)

Index	GARCH	HAR	LSTM	NBEATSx	NHITS	NFM-GNN
SPX	0.394	0.332	0.371	0.331	0.325	0.325
GDAXI	0.330	0.297	0.301	0.285	0.272	0.260
FCHI	0.337	0.322	0.318	0.298	0.284	0.273
FTSE	0.384	0.313	0.322	0.322	0.287	0.290
OMXSPI	0.309	0.246	0.233	0.221	0.214	0.209
N225	0.389	0.283	0.284	0.281	0.272	0.265
KS11	0.252	0.264	0.251	0.228	0.229	0.225
HSI	0.268	0.266	0.245	0.238	0.238	0.228

Table 6.2: Volatility Forecasting Performance (RMSE)

Index	GARCH	HAR	LSTM	NBEATSx	NHITS	NFM-GNN
SPX	0.569	0.560	0.681	0.571	0.540	0.603
GDAXI	0.503	0.458	0.542	0.472	0.449	0.477
FCHI	0.543	0.499	0.599	0.518	0.497	0.515
FTSE	0.560	0.541	0.625	0.557	0.490	0.562
OMXSPI	0.416	0.373	0.410	0.361	0.354	0.379
N225	0.500	0.424	0.480	0.456	0.431	0.444
KS11	0.416	0.415	0.433	0.398	0.399	0.376
HSI	0.400	0.402	0.385	0.381	0.377	0.371

at the 5% significance level.

MAE results reveal that the **NFM-GNN model achieves the lowest error on 7 out of 8 indices**, including SPX, GDAXI, FCHI, OMXSPI, N225, and HSI. This shows the strength of the model to capture the dynamics of average volatility, especially in less stable or nonlinear markets. Compared to traditional econometric models (e.g., HAR) and deep learning baselines (e.g., LSTM, NBEATSx), NFM-GNN offers improved mean-level accuracy, likely due to its frequency-aware representation and integration of temporal graph structures.

The MAE results highlight the overall high accuracy of NFM-GNN, whereas the RMSE results present a more nuanced picture. Although NFM-GNN achieves relatively low error magnitudes across most indices, it records the lowest RMSE value on only **2 of 8 indices**, KS11 and HSI. This outcome does not necessarily indicate inferior performance; rather, it likely reflects the model's sensitivity to a few extreme volatility spikes, such as those observed during the first half of 2020 (see Figure 6.1). These rare but severe fluctuations

can disproportionately inflate the squared-error metric in the short term, even though such deviations tend to revert to stability over time. The RMSE results indicate that NFM-GNN is slightly less robust to extreme forecast errors compared with models such as NHITS, which benefit from explicit multi-scale smoothing and residual interpolation mechanisms. As discussed later in Section 6.4, NFM-GNN still exhibits statistically significant advantages over all baseline models under RMSE criteria. NFM-GNN continues to deliver stable and closely aligned performance across the remaining indices, and subsequent DM test results further confirm its overall superiority and preference for stability in long-term forecasting. The DM test finding suggests that even when penalized by the presence of extreme values, the model maintains strong forecast quality and generalization stability.

The results also show that NFM-GNN consistently remains within the **Model Confidence Set (MCS)** across all stock indices in the forecasting task, demonstrating strong robustness and predictive accuracy. In contrast, most competing models are excluded from the MCS in the majority of cases, indicating inferior or less stable performance.

In summary, **NFM-GNN demonstrates superior MAE performance and competitive RMSE stability**, confirming its effectiveness for volatility forecasting in both mature and emerging equity markets. These findings validate the importance of integrating frequency-domain modeling with graph-based temporal structures.

6.3 Ablation Study

To evaluate the contribution of each architectural component in the proposed NFM-GNN, we conduct a series of controlled ablation experiments under identical data splits and optimization settings. Each variant is compared against the **Full** model across all eight indices.

Normalization. To eliminate scale effects arising from differing volatility levels among indices, we compute the relative change in each error metric (MAE or RMSE) within each

index and then average these values across indices:

$$\Delta_i = \frac{M_i^{(\text{variant})} - M_i^{(\text{Full})}}{M_i^{(\text{Full})}}, \quad (6.1)$$

$$\bar{\Delta} = \frac{1}{N} \sum_{i=1}^N \Delta_i, \quad N = 8. \quad (6.2)$$

Negative $\bar{\Delta}$ values indicate improvement relative to the **Full** model, while positive values reflect deterioration.

Model Variants.

- **Full:** Includes three GNN modules (temporal, feature, and spatial) with masking enabled.
- **w/o Temporal-GNN, w/o Feature-GNN, w/o Spatial-GNN:** remove the respective graph encoder to evaluate its individual contribution.
- **w/o All-GNNs:** replace all GNNs with parameter-matched MLPs to isolate the effect of graph learning.
- **w/o Mask:** disable zero/holiday masking in attention and aggregation layers.
- **w/o Corr:** remove the correlation prior used in the feature graph construction.

Results. Table 6.3 and Table 6.4 present the MAE and RMSE results, respectively, across all eight volatility indices. Table 6.5 reports the macro-averaged relative change ($\bar{\Delta}$) in MAE and RMSE for the 22-day forecasting horizon.

In Tables 6.3 and 6.4, the highest MAE or RMSE value in each row is highlighted in bold. It can be observed that, except for the last index (HSI) where the performance of the original model is suboptimal, our proposed model achieves the best results across the other seven stock indices. Among the ablation variants, the components that exert the most pronounced influence include the mask module used for handling missing values.

Table 6.3: MAE Comparison across Variants

Variant	SPX	GDAXI	FCHI	FTSE	OMXSPI	N225	KS11	HSI
Full (Original)	0.3225	0.2576	0.2711	0.2882	0.2078	0.2631	0.2232	0.2286
w/o Temporal-GNN	0.3242	0.2623	0.2779	0.2948	0.2126	0.2693	0.2278	0.2280
w/o Feature-GNN	0.3264	0.2587	0.2755	0.2907	0.2089	0.2664	0.2261	0.2275
w/o Spatial-GNN	0.3302	0.2615	0.2794	0.2939	0.2126	0.2669	0.2286	0.2280
w/o All-GNNs	0.3262	0.2628	0.2796	0.2932	0.2123	0.2683	0.2256	0.2301
w/o Mask	0.3350	0.2649	0.2794	0.2976	0.2133	0.2676	0.2300	0.2287
w/o Corr	0.3315	0.2632	0.2794	0.2958	0.2121	0.2687	0.2300	0.2285

Table 6.4: RMSE Comparison across Variants

Variant	SPX	GDAXI	FCHI	FTSE	OMXSPI	N225	KS11	HSI
Full (Original)	0.5984	0.4734	0.5115	0.5579	0.3761	0.4406	0.3733	0.3708
w/o Temporal-GNN	0.6024	0.4770	0.5170	0.5625	0.3798	0.4534	0.3800	0.3702
w/o Feature-GNN	0.6011	0.4753	0.5159	0.5594	0.3772	0.4463	0.3787	0.3685
w/o Spatial-GNN	0.6086	0.4793	0.5211	0.5630	0.3807	0.4447	0.3816	0.3695
w/o All-GNNs	0.6030	0.4790	0.5192	0.5622	0.3800	0.4467	0.3777	0.3709
w/o Mask	0.6115	0.4820	0.5191	0.5661	0.3837	0.4489	0.3837	0.3697
w/o Corr	0.6089	0.4813	0.5205	0.5649	0.3810	0.4493	0.3838	0.3683

Each ablation demonstrates its contribution: almost every component, when removed, results in the highest error for at least one index, with the exception of the w/o Feature-GNN variant.

However, Table 6.5 shows that the w/o Feature-GNN configuration exhibits the largest macro-averaged relative change, indicating that feature-level graph learning remains crucial to the overall performance. The model heavily depends on feature-level graph learning to capture inter-index dependencies. After introducing the 5-day and 22-day moving-average (5MA and 22MA) inputs, the network cannot inherently infer how these new features relate to the original volatility sequence without structural guidance. The Feature-GNN explicitly encodes the correlations between the raw volatility series and its short- and long-horizon moving-average representations. Without this relational guidance, the model fails to effectively associate transient fluctuations with their smoothed temporal trends, leading to a considerable decline in predictive accuracy.

The pronounced degradation observed in the w/o Corr variant further supports this interpretation. Because the correlation prior constitutes a core component of the Feature-GNN, its removal weakens the model’s ability to encode pairwise dependencies between the original features and their moving-average extensions. Experimental results indicate

Table 6.5: Ablation Study: Macro-Averaged Relative Change ($\bar{\Delta}$) in MAE and RMSE

Model Variant	MAE $\bar{\Delta}$	RMSE $\bar{\Delta}$
w/o Temporal-GNN	0.0186	0.0122
w/o Feature-GNN	0.0255	0.0167
w/o Spatial-GNN	0.0174	0.0098
w/o All-GNNs	0.0170	0.0110
w/o Mask	0.0085	0.0055
w/o Corr	0.0224	0.0148

that neither using the correlation matrix alone nor relying solely on attention mechanisms can achieve comparable performance. We hypothesize that the correlation prior provides a static, domain-informed relational scaffold that anchors the attention mechanism, while the attention layers dynamically adapt these relationships to current market conditions. Together, they establish a balance between stability (from the prior) and adaptability (from the attention), yielding a more robust representation of feature interactions.

6.4 Statistical Comparison via DM Test and MCS Test

To evaluate the statistical significance of forecasting performance differences, we conduct pairwise Diebold-Mariano (DM) tests comparing NFM-GNN with each baseline model across all stock indices, using both MAE and MSE as evaluation criteria. We summarize in Figure 6.4 the number of indices where NFM-GNN significantly outperforms or underperforms other models at the 5% significance level.

Under the MAE criterion, NFM-GNN achieves widespread dominance. It significantly outperforms LSTM and GARCH on 8 indices each, HAR on 7 indices, NBEATSx on 6 indices and NHITS on 2 indices. Notably, no baseline model shows any statistically significant win over NFM-GNN under MAE, highlighting its consistent accuracy in minimizing absolute error.

Under the MSE criterion, NFM-GNN again demonstrates strong performance, significantly outperforming LSTM on 8 indices, GARCH on 4 indices, and HAR on 2 indices. However, NHITS outperforms NFM-GNN on 2 indices, indicating that while NFM-GNN




Figure 6.4: Significant wins and losses of NFM-GNN relative to baseline models under MAE (left) and RMSE (right) criteria at the 5% significance level based on Diebold–Mariano (DM) tests.

excels at reducing squared errors in most cases, NHITS remains competitive in specific scenarios. NBEATSx shows no significant difference compared to NFM-GNN under RMSE.

Across all indices, traditional econometric approaches (GARCH, HAR) are consistently eliminated from the superior model set ($p\text{-value} \approx 0$), reaffirming their limited ability to capture nonlinear volatility dynamics and cross-market spillovers. Although deep learning baselines (LSTM, NBEATSx, NHITS) occasionally survive the MCS filtering for certain markets (e.g., SPX, FTSE), their inclusion is sporadic and lacks robustness across datasets, highlighting their sensitivity to market-specific structures. In contrast, the proposed **NFM-GNN** persistently attains $p\text{-values}$ close to 1.000 across all markets and both loss functions, remaining in the superior set under every specification. This consistent dominance indicates that NFM-GNN not only achieves state-of-the-art accuracy but also maintains statistical superiority in a way that competing models fail to replicate. Hence, the MCS evidence strongly supports the robustness, stability, and cross-market generalizability of the NFM-GNN framework, while exposing the fragility of conventional and baseline alternatives.

Table 6.6: MCS Test (MAE)

Index	GARCH	HAR	LSTM	NBEATS _x	NHITS	NFM-GNN
SPX	0.000	0.456	0.006	0.456	1.000	0.603
GDAXI	0.000	0.000	0.002	0.002	0.068	1.000
FCHI	0.000	0.000	0.003	0.008	0.066	1.000
FTSE	0.000	0.001	0.003	0.002	1.000	0.481
OMXSPI	0.000	0.000	0.005	0.028	0.186	1.000
N225	0.000	0.154	0.089	0.154	0.154	1.000
KS11	0.001	0.000	0.013	0.383	0.383	1.000
HSI	0.000	0.000	0.000	0.008	0.008	1.000

Table 6.7: MCS Test (RMSE)

Index	GARCH	HAR	LSTM	NBEATS _x	NHITS	NFM-GNN
SPX	0.000	0.449	0.007	0.449	1.000	0.613
GDAXI	0.000	0.000	0.000	0.001	0.054	1.000
FCHI	0.000	0.000	0.004	0.007	0.068	1.000
FTSE	0.000	0.001	0.007	0.007	1.000	0.453
OMXSPI	0.000	0.000	0.002	0.038	0.166	1.000
N225	0.000	0.146	0.088	0.146	0.146	1.000
KS11	0.000	0.000	0.014	0.424	0.424	1.000
HSI	0.000	0.000	0.000	0.002	0.005	1.000

Collectively, the DM and MCS results provide compelling statistical evidence that the NFM-GNN framework consistently outperforms conventional and neural baselines across diverse markets. Its ability to remain within the superior model set under both MAE and RMSE criteria, while maintaining dominance in pairwise significance tests, demonstrates not only predictive accuracy but also cross-market adaptability and methodological robustness.

These findings establish NFM-GNN as a statistically validated advancement in volatility forecasting, capable of delivering reliable performance under varying market dynamics.

Chapter 7

Discussion

Given the increasing interdependence of global financial systems, reliable cross-market volatility forecasts are critical for systemic risk surveillance and macro-prudential decision-making. This study set out to examine whether the proposed NFM-GNN can enhance volatility forecasting across multiple financial markets. The empirical results demonstrate that NFM-GNN consistently achieves lower forecasting errors than all benchmark models, with statistically significant improvements confirmed by both the Diebold–Mariano (DM) and Model Confidence Set (MCS) tests. These findings indicate that integrating multi-graph-based learning mechanisms yields measurable gains in predictive accuracy and stability across heterogeneous markets. By combining spectral decomposition and relational graph learning, the model effectively captures the multi-scale and cross-market structures that underlie volatility dynamics, an ability that conventional econometric and sequential deep-learning frameworks largely lack.

7.1 Interpretation of the Main Findings

The superior performance of NFM-GNN can be attributed to its triple representation of temporal, feature and cross-sectional dependencies. The module encodes long- and short-horizon dynamics (e.g., 5-day and 22-day moving averages), which helps the model bal-

ance responsiveness to short-term fluctuations with awareness of long-term persistence in realized volatility. Meanwhile, the multi-graph component learns structural relations among indices, enabling the model to capture volatility spillovers, contagion channels, and systemic linkages across global markets. This joint modeling strategy allows NFM-GNN to remain stable under high market turbulence—such as during the COVID-19 shock—when volatility transmission intensifies and traditional single-series approaches typically fail. The consistent DM and MCS superiority across eight indices suggests that the improvement is not data-specific but structurally driven

From a theoretical standpoint, these findings suggest that volatility is not merely an autoregressive outcome but a networked phenomenon emerging from interdependent market behavior. The ability to represent such interactions through graph structures provides a new computational lens for understanding financial connectedness and the propagation of uncertainty. This contributes to the growing body of literature that emphasizes the importance of network-based and multi-scale representations in financial econometrics.

7.2 Comparison with Existing Studies

The results extend existing research in several meaningful ways. Compared with classical models such as GARCH and HAR, which rely on static assumptions and short-memory processes, NFM-GNN flexibly captures nonlinearities and long-range temporal dependencies. In contrast to neural networks such as LSTM or architectures such as NBEATSx and NHITS, the proposed model embeds structural awareness through multi-graph learning. This integration of frequency-domain modeling and graph representation situates NFM-GNN at the intersection of econometric theory and modern machine learning. It also demonstrates that relational learning is crucial for achieving robust generalization across markets with diverse dynamics and data-generating processes.

7.3 Practical and Economic Implications

From a practical perspective, the model’s robustness across markets makes it a valuable tool for a wide range of financial stakeholders. Portfolio managers can use the forecasts for dynamic risk budgeting, option pricing, and cross-market hedging strategies. Beyond direct forecasting use, the framework contributes to the construction of early-warning indicators for systemic fragility, enabling data-driven stress-testing scenarios for financial regulators. Regulatory bodies and central banks may employ the model’s network-based outputs to monitor systemic risk propagation, identify volatility hubs, and evaluate the resilience of financial systems to global shocks.

The graph-based component also enhances interpretability: it allows analysts to visualize how shocks originating in one market (e.g., U.S. equities) influence volatility in others (e.g., Europe or Asia), providing a transparent analytical basis for macroprudential policy. Moreover, because the model’s structure naturally accommodates multiple time horizons, it offers strong potential for extension into high-frequency applications, where rapid volatility estimation is critical for both traders and financial regulators.

7.4 Limitations and Data Challenges

Despite its strong performance, several limitations merit consideration.

First, the analysis is conducted on daily realized volatility, which, while widely available and reliable, may not fully reflect intraday market behavior. High-frequency realized volatility (e.g., 5-minute intervals) provides richer information but is typically proprietary, noisy, and requires rigorous filtering to mitigate microstructure biases and asynchronous trading effects.

Second, the study assumes a fixed inter-market graph, whereas real financial networks evolve over time as correlations shift with changing macroeconomic regimes.

Third, the model’s complexity introduces computational overhead, which may constrain real-time implementation in ultra-high-frequency environments.

These challenges highlight the inherent trade-off between model expressiveness and operational feasibility in deploying deep learning systems in real-time financial environments. They do not detract from the framework’s core contribution but rather highlight avenues for refinement and extension.

7.5 Directions for Future Research

Future work may build upon this study in several directions.

Methodologically, incorporating magnetic Laplacian graph neural networks could enable the model to learn directional and asymmetric dependencies—features that are crucial for representing real-world volatility spillovers and lead–lag effects. X. Zhang, Y. He, Brugnone, Perlmutter, and Hirn (2021) introduced MagNet, a GNN designed for directed graphs using a complex Hermitian operator known as the magnetic Laplacian, which encodes geometric structure in its magnitude and directional information in its phase. By tuning a charge parameter to capture variations across directed cycles, MagNet effectively models asymmetric relationships and has shown superior performance across node classification and link prediction benchmarks. Adapting these principles to volatility forecasting would allow the network to better represent asymmetric information flow and feedback among financial assets.

Extending complex-valued neural architectures also presents a valuable avenue. Xavier and John (2025) demonstrated the advantages of complex-valued spatio–temporal graph convolutional networks (CSGCNNs) in integrating multimodal information by modelling amplitude–phase interactions, highlighting the potential of complex representations for handling phase relationships and multi-tensor fusion. Applying similar complex-valued graph mechanisms to volatility forecasting could therefore enhance the model’s ability

to encode amplitude–phase interactions and temporal correlations inherent in market dynamics.

A further direction involves embedding graph learning directly in the frequency domain, allowing spectral filters to operate over graph topologies and thereby improving stability and representation in complex-valued GNN formulations. From a data perspective, incorporating additional market information such as open-to-close price spreads, trading volumes, or macroeconomic indicators could enrich the feature space and improve predictive stability during regime shifts. Finally, adopting dynamic or attention-based graph mechanisms would allow the network to adaptively update its structure in response to evolving market connectivity, enhancing both interpretability and real-time adaptability.

7.6 Integrative Reflection

Taken together, these findings demonstrate that NFM-GNN is not merely a technical enhancement but represents a conceptual step toward modeling financial volatility as a dynamic, interconnected system.

By embedding graph relations within a frequency-aware neural architecture, the model bridges the gap between traditional econometric intuition and the expressive capacity of modern deep learning.

Its ability to outperform state-of-the-art alternatives across multiple markets underscores both its empirical robustness and its potential as a foundational framework for future research.

As high-frequency data become increasingly available, and as financial systems continue to evolve in complexity, the principles underlying NFM-GNN offer a promising blueprint for the next generation of volatility forecasting models.

References

- Andersen, Torben G and Tim Bollerslev (1998). “Answering the skeptics: Yes, standard volatility models do provide accurate forecasts”. In: *International Economic Review*, pp. 885–905.
- Andersen, Torben G, Tim Bollerslev, Peter Christoffersen, and Francis X Diebold (2007). “Practical volatility and correlation modeling for financial market risk management”. In: *The Risks of Financial Institutions*. University of Chicago Press, pp. 513–548.
- Audrino, Francesco and Simon D Knaus (2016). “Lassoing the HAR model: A model selection perspective on realized volatility dynamics”. In: *Econometric Reviews* 35.8-10, pp. 1485–1521.
- Barndorff-Nielsen, Ole E and Neil Shephard (2002). “Econometric analysis of realized volatility and its use in estimating stochastic volatility models”. In: *Journal of the Royal Statistical Society Series B: Statistical Methodology* 64.2, pp. 253–280.
- Bollerslev, Tim (1986). “Generalized autoregressive conditional heteroskedasticity”. In: *Journal of Econometrics* 31.3, pp. 307–327.
- Breiman, Leo (2001). “Random forests”. In: *Machine Learning* 45.1, pp. 5–32.
- Challu, Cristian, SKin G. Olivares, Boris N. Oreshkin, Federico Garza Ramirez, Max Mergenthaler Canseco, and Artur Dubrawski (2023). “NHITS: Neural Hierarchical Interpolation for Time Series Forecasting”. In: *The Association for the Advancement of Artificial Intelligence*. Vol. 37. 6, pp. 6989–6997.
- Chen, Luyang, Markus Pelger, and Jason Zhu (2024). “Deep learning in asset pricing”. In: *Management Science* 70.2, pp. 714–750.

- Chen, Tianqi and Carlos Guestrin (2016). “XGBoost: A scalable tree boosting system”. In: *Proceedings of the 22nd Association of Computing Machinery’s Special Interest Group for Knowledge Discovery and Data Mining*, pp. 785–794.
- Cheng, Rui and Qing Li (2021). “Modeling the momentum spillover effect for stock prediction via attribute-driven graph attention networks”. In: *Proceedings of the AAAI conference on artificial intelligence*. Vol. 35. 1, pp. 55–62.
- Chi, Zhengyang, Junbin Gao, and Chao Wang (2024). “Global Stock Market Volatility Forecasting Incorporating Dynamic Graphs and All Trading Days”. In: *arXiv:2409.15320*.
- Christensen, Kim, Mathias Siggaard, and Bezirgen Veliyev (2023). “A machine learning approach to volatility forecasting”. In: *Journal of Financial Econometrics* 21.5, pp. 1680–1727.
- Cont, Rama (2001). “Empirical properties of asset returns: stylized facts and statistical issues”. In: *Quantitative finance* 1.2, p. 223.
- Corsi, Fulvio (2009). “A Simple Approximate Long-Memory Model of Realized Volatility”. In: *Journal of Financial Econometrics* 7.2, pp. 174–196.
- Corsi, Fulvio, Francesco Audrino, and Roberto Renò (2012). “HAR Modeling for Realized Volatility Forecasting”. In: *Handbook of Volatility Models and Their Applications*. Hoboken, NJ, USA: John Wiley & Sons, Inc., pp. 363–382. ISBN: 9780470872512.
- Dhingra, Barkha, Shallu Batra, Vaibhav Aggarwal, Mahender Yadav, and Pankaj Kumar (2024). “Stock market volatility: a systematic review”. In: *Journal of Modelling in Management* 19.3, pp. 925–952.
- Engle, Robert F. (1982). “Autoregressive Conditional Heteroskedasticity with Estimates of the Variance of United Kingdom Inflation”. In: *Econometrica* 50.4, pp. 987–1008.
- Fama, Eugene F (1991). “Efficient capital markets: II”. In: *The journal of finance* 46.5, pp. 1575–1617.
- Feng, Fuli, Xiangnan He, Xiang Wang, Cheng Luo, Yiqun Liu, and Tat-Seng Chua (2019). “Temporal Relational Ranking for Stock Prediction”. In: *Association for Computing Machinery Transactions on Information Systems* 37.2, pp. 1–30.

- Glosten, Lawrence R., Ravi Jagannathan, and David E. Runkle (1993). “On the Relation Between the Expected Value and the Volatility of the Nominal Excess Return on Stocks”. In: *The Journal of Finance* 48.5, pp. 1779–1801.
- Hochreiter, Sepp and Jürgen Schmidhuber (1997). “Long Short-Term Memory”. In: *Neural Computation* 9.8, pp. 1735–1780.
- Hu, Zinan and Sumuya Borjigin (2024). “The amplifying role of geopolitical Risks, economic policy Uncertainty, and climate risks on Energy-Stock market volatility spillover across economic cycles”. In: *The North American Journal of Economics and Finance* 71, p. 102114.
- Jin, Jie, Qian Su, Yan Xu, Zhengrui He, and Yu Lu (2022). “Efficient Radiation Pattern Prediction of Array Antennas Based on Complex-Valued Graph Neural Networks”. In: *Institute of Electrical and Electronics Engineers Antennas and Wireless Propagation Letters* 21.12, pp. 2467–2471.
- Kim, Minjung, Yusuke Hioka, and Michael Witbrock (2024). “Neural Fourier Modelling: A Highly Compact Approach to Time-Series Analysis”. In: *arXiv:2410.04703*.
- Kürüm, Efsun, Gerhard-Wilhelm Weber, and Cem Iyigun (2018). “Early warning on stock market bubbles via methods of optimization, clustering and inverse problems”. In: *Annals of Operations Research* 260.1, pp. 293–320.
- Li, Yaguang, Rose Yu, Cyrus Shahabi, and Yan Liu (2018). “Diffusion Convolutional Recurrent Neural Network: Data-Driven Traffic Forecasting”. In: *International Conference on Learning Representations*.
- Li, Yanan and David E Giles (2015). “Modelling volatility spillover effects between developed stock markets and Asian emerging stock markets”. In: *International Journal of Finance & Economics* 20.2, pp. 155–177.
- Luong, Chuong and Nikolai Dokuchaev (2018). “Forecasting of realised volatility with the random forests algorithm”. In: *Journal of Risk and Financial Management* 11.4, p. 61.
- Mehtab, Sidra, Jaydip Sen, and Subhasis Dasgupta (2020). “Robust Analysis of Stock Price Time Series Using CNN and LSTM-Based Deep Learning Models”. In: *Insti-*

tute of Electrical and Electronics Engineers International Conference on Electrical, Computer and Communication Technologies.

- Merton, Robert C (1973). “An intertemporal capital asset pricing model”. In: *Econometrica: Journal of the Econometric Society*, pp. 867–887.
- Müller, Ulrich A, Michel M Dacorogna, Rakhal D Davé, Richard B Olsen, Olivier V Pictet, and Jacob E Von Weizsäcker (1997). “Volatilities of different time resolutions—analyzing the dynamics of market components”. In: *Journal of Empirical Finance* 4.2-3, pp. 213–239.
- Nelson, Daniel B. (1991). “Conditional Heteroskedasticity in Asset Returns: A New Approach”. In: *Econometrica* 59.2, pp. 347–370.
- Nelson, David M.Q., Adriano C. M. Pereira, and Renato A de Oliveira (2017). “Stock market’s price movement prediction with LSTM neural networks”. In: *International Joint Conference on Neural Networks*, pp. 2161–4407.
- Olivares, Kin G., Cristian Challu, Grzegorz Marcjasz, Rafał Weron, and Artur Dubrawski (2023). “Neural basis expansion analysis with exogenous variables: Forecasting electricity prices with NBEATSx”. In: *International Journal of Forecasting* 39.2, pp. 884–900.
- Oreshkin, Boris N, Dmitri Carpov, Nicolas Chapados, and Yoshua Bengio (2019). “N-BEATS: Neural basis expansion analysis for interpretable time series forecasting”. In: *arXiv:1905.10437*.
- Poon, Ser-Huang and Clive W J Granger (2003). “Forecasting volatility in financial markets: A review”. In: *Journal of Economic Literature* 41.2, pp. 478–539.
- Sharpe, William F (1964). “Capital asset prices: A theory of market equilibrium under conditions of risk”. In: *The journal of finance* 19.3, pp. 425–442.
- Sitzmann, Vincent, Julien Martel, Alexander Bergman, David Lindell, and Gordon Wetzstein (2020). “Implicit Neural Representations with Periodic Activation Functions”. In: *Conference on Neural Information Processing Systems*.
- Son, Bumho, Yunyoung Lee, Seongwan Park, and Jaewook Lee (2023). “Forecasting global stock market volatility: The impact of volatility spillover index in spatial-temporal graph-based model”. In: *Journal of Forecasting*.

- Souto, Hugo Gobato (2025). “Evaluating the Efficacy of NHITS for Forecasting Stock Realized Volatility: A Comparative Analysis with Established Models”. In: *Computational Economics*.
- Souto, Hugo Gobato and Amir Moradi (2024). “Introducing NBEATSx to realized volatility forecasting”. In: *Expert Systems with Applications* 242, p. 122802.
- Tsai, I-Chun (2014). “Spillover of fear: Evidence from the stock markets of five developed countries”. In: *International Review of Financial Analysis* 33, pp. 281–288.
- Vapnik, Vladimir (1995). *The Nature of Statistical Learning Theory*. Springer New York.
- Wu, Haixu, Jiehui Xu, Jianmin Wang, and Mingsheng Long (2021). “Autoformer: Decomposition transformers with auto-correlation for long-term series forecasting”. In: *Conference on Neural Information Processing Systems* 34, pp. 22419–22430.
- Wu, Tong, Anna Scaglione, and Daniel Arnold (2024). “Complex-Value Spatiotemporal Graph Convolutional Neural Networks and Its Applications to Electric Power Systems AI”. In: *Institute of Electrical and Electronics Engineers Transactions on Smart Grid* 15.3, pp. 3193–3207.
- Xavier, S. P. and S. P. John (2025). “Advanced emotion recognition from facial images and speech signals using complex-value spatio-temporal graph convolutional neural networks”. In: *Signal, Image and Video Processing* 19, p. 939.
- Xu, Ke, You Wu, Haohao Xia, Ningjing Sang, and Bingxing Wang (2022). “Graph Neural Networks in Financial Markets: Modeling Volatility and Assessing Value-at-Risk”. In: *Journal of Computer Technology and Software*.
- Yi, Kun, Qi Zhang, Wei Fan, Hui He, Liang Hu, Pengyang Wang, Ning An, Longbing Cao, and Zhendong Niu (2023). “FourierGNN: Rethinking multivariate time series forecasting from a pure graph perspective”. In: *Advances in Neural Information Processing Systems* 36, pp. 69638–69660.
- Yi, Qian, Lin Xiao, Yongsheng Zhang, Bolin Liao, Lei Ding, and Hua Peng (2018). “Non-linearly Activated Complex-Valued Gradient Neural Network for Complex Matrix Inversion”. In: *2018 Ninth International Conference on Intelligent Control and Information Processing*, pp. 44–48.

- Yu, Bing, Haoteng Yin, and Zhanxing Zhu (2018). “Spatio-Temporal Graph Convolutional Networks: A Deep Learning Framework for Traffic Forecasting”. In: *International Joint Conferences on Artificial Intelligence*, pp. 3634–3640.
- Zhang, Xitong, Yixuan He, Nathan Brugnone, Michael Perlmutter, and Matthew Hirn (2021). “MagNet: A Neural Network for Directed Graphs”. In: *Advances in Neural Information Processing Systems*. Ed. by M. Ranzato, A. Beygelzimer, Y. Dauphin, P.S. Liang, and J. Wortman Vaughan. Vol. 34. Curran Associates, Inc., pp. 27003–27015.
- Zhong, Yi and Jiapeng Liu (2021). “Correlations and volatility spillovers between China and Southeast Asian stock markets”. In: *The Quarterly Review of Economics and Finance* 81, pp. 57–69.
- Zhou, Tian, Ziqing Ma, Qingsong Wen, Xue Wang, Liang Sun, and Rong Jin (2022). “FEDformer: Frequency Enhanced Decomposed Transformer for Long-term Series Forecasting”. In: *International Conference on Machine Learning*.